

GDX Factor-Residual Reversal Strategy

Research summary: what the code does, what it found, and what to do next

In-sample data 2006-05-22 → 2021-12-31 · Frozen 2026-09-22 at commit `85093c09a2` · repository `sreejaghose/GDX-Factor-Model`, branch `claude/charming-albattani-a06ljy`

1. Executive summary

Bottom line. The pipeline is complete, tested (83 automated tests) and reproducible, and it applied a strict, pre-declared selection protocol to 10,800 strategy configurations. **No configuration met all of the stated requirements on the validation period (2017-2021).** The strategy that was frozen is the best available by the pre-declared tie-break rule and carries the status flag `no_candidate_passed_validation_filters`. The evidence points to a real but thin day-2 reversal effect in GDX's factor residual that is not strong enough, after correcting for the number of configurations tried, to be called a robust edge.

- **Idea.** Explain GDX's daily excess return with SPY/GLD/TLT (rolling OLS). When GDX moves much more or less than its factor exposures imply (a large standardised residual z), bet on partial reversal over the next H days.
- **Data reproduced exactly.** The Python pipeline reproduces the workbook's Full-Sample Sweep for all 15 factor subsets (max difference $\approx 1e-10$; SPY+GLD: $\alpha -3.93$ bps/day, $\beta_{\text{SPY}} 0.498$, $\beta_{\text{GLD}} 1.744$, $R^2 0.6308$).
- **Search.** 10,800 configurations (3 factor sets $\times$ 5 lookbacks $\times$ 4 accumulation windows $\times$ 6 thresholds $\times$ 5 holding periods $\times$ 3 position policies $\times$ 2 direction modes), ranked on 2008-05-27 → 2016-12-30 (2,167 trading days), net of 2 bps.
- **Funnel.** Only 10 of 10,800 configurations passed all six development hard filters; all 10 failed at least one filter on validation (every one failed the "still profitable without the best 1% of days" test).
- **Frozen primary:** `GLD|L20|m1|k0.75|H2|A|rev` — GLD-only factor model, 20-day lookback, 1-day residual, $|z| \geq 0.75$, hold 2 days, run-out policy, imposed mean reversion. Net Sharpe **0.83** (dev) → **0.46** (validation); gross 0.90 → 0.54; ~ 6.4 new trades per month; trades every month in both periods.
- **Multiple-testing reality check.** Deflated Sharpe Ratio of the primary = **0.041** (needs ≥ 0.95): with 10,800 trials, a dev Sharpe of about 1.41 would be expected from luck alone. Yet it beats 99.6% of random-timing placebos, and inverting it loses (Sharpe -0.96) — so the timing contains information, just not enough.
- **Walk-forward.** Re-applying the selection rule each year 2011-2021 picked 9 different configurations in 11 years and produced a stitched out-of-sample Sharpe of **-0.19** — the selection process itself is unstable.
- **Key diagnostic.** The edge sits on **day 2** after the signal, not day 1; the Stage-2 one-day forecast therefore cannot see it (for the primary's settings the estimated slope $\hat{y}$ even points the wrong way), which is why every shortlisted config uses the imposed "reversion" direction.

2. The task and its requirements

Use the GDX factor model in `gold_etf_factor_regression.xlsx` to forecast GDX's next-day return, turn the forecast into a long / short / flat decision, back-test a grid of candidate strategies in-sample, select a small set of robust candidates, freeze them, and ship a runner that applies the frozen strategy to out-of-sample data later with no re-optimisation.

Requirements:

- At least one new trade per calendar month.
- Robust to parameter changes — good results must hold in the neighbourhood of a configuration, not at an isolated peak.
- Consistent over time — not driven by one year or a handful of days.
- Signal observed at the close of day t using prices as of that close, and executed at that same close.
- Holding period (exit time, H days) is a tunable parameter like lookback and threshold.

3. Data

- **Prices (input):** daily closes for SPY, IAU, GLD, TLT, GDX, GDXJ; 3,932 rows, 2006-05-22 → 2021-12-31. Zero or blank prices are treated as missing; GDXJ is dropped (not needed; its pre-inception blanks were stored as zeros).
- **FedFunds (input):** monthly effective Fed Funds rate (FRED FEDFUNDS). Daily risk-free = FEDFUNDS/100/252.
- **Returns:** simple close-to-close returns; excess return = return − RF/252. After the first (return-less) row: 3,931 days.
- **Risk-free look-ahead fix.** The workbook maps each day to *that same month's* FEDFUNDS average, which is not known until the month ends. The code implements `rf_mode = same_month` (only to reproduce Excel) and `prior_month` (used everywhere in the strategy). The difference is at most 0.38 bps/day.
- **Adjusted closes.** A heuristic ex-dividend test found no significant dividend-day drops (SPY −20.7 bps vs −45 expected if unadjusted, $t = -1.52$; TLT −2.9 bps vs −25, $t = -0.38$), and SPY's 2006 level (87.15 vs ~126 actual) confirms back-adjustment.
- **Reproduction of the workbook.** Full-sample OLS on every one of the 15 factor subsets matches the Full-Sample Sweep sheet (alpha, betas, t-stats, p-values, R^2 , adj. R^2 , AIC, BIC) within $1e-10$, well inside the $1e-4$ tolerance. Best BIC: SPY+GLD (−21,304.2); best AIC: SPY+GLD+TLT (β_{TLT} 0.085, t 2.69). IAU and GLD are never used together.

4. What the code does — the model

4.1 Stage 1: rolling factor residual (`src/factor_model.py`)

For factor set $F \in \{\text{GLD}, \text{SPY}+\text{GLD}, \text{SPY}+\text{GLD}+\text{TLT}\}$ and lookback L , at each day t an OLS with intercept is fitted on the window $[t-L, t-1]$ — day t is excluded so today's move cannot contaminate its own benchmark. The out-of-window residual is $e_t = \text{ExRet}_{\text{GDX},t} - \hat{\alpha} - \hat{\beta}^T \cdot \text{ExRet}_{F,t}$, the residual volatility $\hat{\sigma}_e$ is the regression standard error $\sqrt{(\text{SSR}/(n-p))}$, and the score is $z_t = \sum_{j < m} e_{t-j} / (\hat{\sigma}_e \sqrt{m})$ for accumulation window m . Rolling sums of $X'X$ and $X'y$ make every window a subtraction with batched solves (≈ 5 ms per (F, L)); results match a brute-force per-day OLS to $1e-8$.

4.2 Stage 2: one-day-ahead forecast (`src/forecast.py`)

At day t , pairs $(z_s, \text{ExRet}_{\text{GDX},s+1})$ for $s \leq t-1$ (expanding window, minimum 250 pairs) are used to fit $\text{ExRet}_{s+1} = \gamma \cdot z_s$ without an intercept (so GDX's drift doesn't create a permanent short bias). The newest target is day t 's own return, known at t 's close. Forecast $\hat{r}_{t+1} = \hat{\gamma}_t \cdot z_t$. The intercept model is kept as a diagnostic, and a rolling $W = 500$ window was run as a sensitivity check (less stable: $\hat{\gamma}$ negative on 65–89% of days vs $\sim 95\%$ expanding).

4.3 Trade rule and position engine (`src/backtest.py`)

- **Entry:** at close t , if $|z_t| \geq k$, take a position in direction $\text{sign}(\hat{r}_{t+1})$ (*estimated* mode) or $-\text{sign}(z_t)$ (*reversion* mode, imposes mean reversion).
- **Position** is always in $\{-1, 0, +1\}$; held from close t to close $t+H$. Policies: **A** run-out (ignore new signals while in a trade), **B** reset/flip (same-direction signal restarts the clock; opposite flips), **C** reset/flat (opposite closes to flat), and **D** stack-to-2x (reference only, never selectable).
- **Timing:** pos_t is the position after the close of t ; strategy return $\text{strat}_{t+1} = \text{pos}_t \cdot \text{ExRet}_{\text{GDX},t+1} - \text{cost}_{t+1}$; $\text{cost} = 2$ bps per unit of position change (entry 1, exit 1, flip 2). Returns are a self-financing 1x overlay in excess of cash.
- The path-dependent loop is compiled with numba (≈ 0.13 ms per back-test) and unit-tested on hand-built signal sequences for every policy.

5. What the code does — the research protocol

- **Grid** (`config/grid.yaml`): factors {GLD, SPY+GLD, SPY+GLD+TLT} $\times$ L {20, 40, 60, 120, 250} $\times$ m {1, 2, 3, 5} $\times$ k {0.5, 0.75, 1.0, 1.25, 1.5, 2.0} $\times$ H {1, 2, 3, 5, 10} $\times$ policy {A, B, C} $\times$ direction {estimated, reversion} = **10,800** configurations (plus 3,600 D-policy references). The full grid runs in ≈ 4 s; every configuration's result — losers included — is persisted.
- **One evaluation window for all**: signals are zeroed until the slowest configuration (L = 250, m = 5 plus 250 Stage-2 pairs) has a valid forecast: 2008-05-23. Scoring starts 2008-05-27.
- **Splits** (`config/splits.yaml`): development 2008-05-27 $\rightarrow$ 2016-12-30 (2,167 days, 103 months) for the full grid; validation 2017-01-03 $\rightarrow$ 2021-12-31 (1,259 days) for the shortlist only. A logged *validation gate* records every look at validation data and refuses a second, different shortlist unless a reason is logged as a protocol deviation.
- **Hard filters** (development, net of 2 bps): trades in 100% of calendar months; ≥ 150 entries (scaled pro rata for shorter windows); net Sharpe > 0 in both halves; positive net return in $\geq 70\%$ of calendar years; still profitable after removing the best 1% of days; policy $\neq$ D.
- **Robustness score**: for each configuration, the median development net Sharpe over its neighbourhood (same factors / policy / direction; each of L, m, k, H at its own value or ± 1 grid step — up to 81 configurations). Configurations are ranked by this median, not their own Sharpe, so isolated peaks score low. Ties: fewer factors, shorter H, policy B.
- **Statistical checks on the top 10**: Deflated Sharpe Ratio (Bailey & López de Prado) with N = 10,800; 1,000 random-entry placebos with the same trade count, holding lengths and long/short mix; sign-flip; stationary block bootstrap (2,000 reps, mean block 10 days) for a 95% CI. All randomness is seeded.
- **Final pick**: evaluate the 10 on validation once; choose a primary and up to 2 alternates that pass the hard filters on validation, by combined rank of validation Sharpe and neighbourhood score. If none passes, freeze the best by fewest failed filters and flag it.
- **Walk-forward**: for each year Y = 2011...2021, apply the same selection rule to data up to Y-1 and trade the choice in Y.
- **Freeze**: `results/frozen_config.json` (parameters, dev and validation metrics, checks, git hash, data hash) plus a recorded SHA-256 (6981dc958cd949d4...) that the OOS runner verifies before running.

5.1 Code map and tests


Module	Role
<code>src/data.py</code>	load prices / Fed Funds, returns, rf_mode, adjusted-close check, IS+OOS concatenation with seam checks
<code>src/factor_model.py</code>	full-sample OLS (Excel reproduction), rolling out-of-window residuals, z-scores, caching
<code>src/forecast.py</code>	Stage-2 predictive regression (expanding / rolling), direction modes
<code>src/backtest.py</code>	numba position engine (policies A-D), P&L with costs, trade list
<code>src/grid.py</code>	grid enumeration, common start date, running configs, persistence
<code>src/metrics.py</code>	performance, frequency, consistency, exposure metrics — vectorised over configs
<code>src/splits.py</code>	dev/validation masks, logged validation gate, anchored walk-forward
<code>src/robustness.py</code>	hard filters, neighbourhood score, DSR, placebo, sign-flip, bootstrap, selection rule
<code>src/report.py</code>	event study, heatmaps, $\hat{\gamma}$ and beta charts, equity, monthly entries, HTML report
<code>scripts/run_is.py</code>	full in-sample workflow: grid $\rightarrow$ filters $\rightarrow$ shortlist $\rightarrow$ validation $\rightarrow$ freeze $\rightarrow$ walk-forward $\rightarrow$ report
<code>scripts/run_oos.py</code>	applies the frozen config to OOS data (hash-verified, no re-optimisation)

83 tests cover: Excel reproduction (15 subsets); brute-force OLS equivalence; Stage-2 vs statsmodels; every position policy on hand-built sequences; cost accounting; metrics vs pandas/statsmodels references; the grid vs single back-tests; a dev-only grid run equal to the dev slice of the full run; the validation gate; walk-forward using only past data; the OOS runner's refusals and seam checks; and the look-ahead checklist — randomising every real price after a date T leaves every z_t , $\hat{r}_{t+1}$ and pos_t for $t \leq T$ unchanged, and a "cheating" oracle on r_{t+1} scores Sharpe > 20 while the same oracle on r_t does not.

6. Findings

6.1 Stage 1 — the factor model and its stability


Rolling betas are far from constant: β_{GLD} ranges $\approx 1.0 \rightarrow 3.5$ (peaking 2016), β_{SPY} falls to ≈ 0 in 2019–20, β_{TLT} hovers near zero. This supports refitting on a rolling window. The residual scores are roughly standard normal (SPY+GLD, $L = 250$: std of z 1.11 for $m = 1$, 0.99 for $m = 5$; $|z| > 2$ on 5–6% of days).



Stage-1 rolling betas for the primary's factor set (GLD), several lookbacks.

6.2 Stage 2 — does the residual predict tomorrow?

Weakly, and not in the way the design assumed. With SPY+GLD, $L = 250$, the full-IS slope $\hat{\gamma}$ was negative (reversal) for $m = 1 \dots 5$, but significant only for $m = 2$ ($t \approx -4.2$). Decomposing by lag showed why: **yesterday's** residual predicts the next day's return ($t = -5.2$), today's barely does ($t = -0.95$), and residuals 2–3 days back predict a small bounce the other way. Raw GDX shows the same pattern (autocorrelation -0.021 at lag 1, -0.069 at lag 2). For the primary's own settings (GLD, $L = 20$, $m = 1$), $\hat{\gamma}$ is *positive* throughout — so the "estimated" direction would bet on continuation, while the profitable trade is reversion two days out.



$\hat{\gamma}_t$ and its t-stat over time (GLD, $L = 20$), development shaded blue, validation orange. The t-stat never exceeds ± 2 after 2009.


6.3 Event study — where the edge sits in time

Average cumulative GDX excess return after signals with $|z| \geq k$, in the trade's direction (bps):

		day 1	day 2	day 3	day 5	day 10
period	k					
development	0.50	-0.4	11.8	6.8	-12.9	-39.2
	0.75	2.0	26.5	19.5	4.4	-9.3
	1.00	-2.1	27.4	22.9	6.5	-19.9
	1.25	-0.7	35.2	29.8	26.0	-1.7
	1.50	-17.6	14.9	12.4	32.6	17.6
	2.00	-9.6	30.3	26.1	90.8	70.5
validation	0.50	6.2	5.6	2.2	-17.9	-2.2
	0.75	12.3	12.5	7.6	-13.9	6.3
	1.00	9.3	12.9	2.1	-13.6	2.1
	1.25	12.2	29.9	19.6	1.3	28.3
	1.50	18.7	31.6	14.5	-15.9	-6.5
	2.00	19.6	62.7	46.3	13.7	71.2

The day-1 concentration hypothesis is not supported in development: day 1 is ≈ 0 (−18 to +2 bps), the edge appears on **day 2** (+12 to +35 bps) and decays by day 4–10. In validation, day 1 carries roughly half the day-2 total. This is the empirical reason $H = 2$ is chosen, and the reason a one-day-ahead forecast is aimed at the wrong day.

Event study, GLD $L=20$ $m=1$, direction=reversion (bands ± 1 s.e.; overlapping events)



Event study for GLD, $L = 20$, $m = 1$, reversion direction; bands are ± 1 s.e. (too narrow — events overlap).

6.4 The whole grid


Distribution of net Sharpe (2 bps) across all 10,800 selectable configurations:

	5th pct	25th pct	median	75th pct	95th pct	max	share > 0
development	-0.59	-0.23	0.02	0.27	0.61	1.20	0.52
validation	-0.83	-0.52	-0.27	0.01	0.37	1.30	0.26
full	-0.53	-0.27	-0.07	0.14	0.42	0.72	0.41

The median configuration is roughly break-even in development and **loses in validation** (median −0.27; only 26% of configs positive). Median Sharpe by parameter value shows what "worked" in development did not carry over:

period	development	validation
direction_mode = estimated	-0.090	-0.326
direction_mode = reversion	0.117	-0.217
m = 1	-0.089	0.010
m = 2	0.254	-0.287
m = 3	0.066	-0.235
m = 5	-0.153	-0.478
H = 1	0.119	-0.075
H = 2	0.175	-0.233
H = 3	0.040	-0.344
H = 5	-0.039	-0.407
H = 10	-0.128	-0.235
policy = A_run_out	-0.013	-0.252
policy = B_reset_flip	0.058	-0.274
policy = C_reset_flat	0.020	-0.286
factors = GLD	0.076	-0.322
factors = SPY+GLD	0.017	-0.242
factors = SPY+GLD+TLT	-0.033	-0.238
lookback = 20	-0.037	-0.193
lookback = 40	-0.061	-0.241
lookback = 60	-0.008	-0.270
lookback = 120	0.084	-0.253
lookback = 250	0.112	-0.358
k = 0.5	-0.029	-0.454
k = 0.75	-0.033	-0.373
k = 1.0	0.029	-0.265
k = 1.25	0.075	-0.211
k = 1.5	0.101	-0.227
k = 2.0	0.020	-0.124


Reversion beat estimated direction in both periods; m = 2 was the best accumulation window in development but among the worst in validation, where only m = 1 is non-negative; short holding periods (H = 1–2) were best in development.




Left: development vs validation Sharpe for every configuration (Spearman rank correlation 0.20). Right: mean validation Sharpe by development-Sharpe decile — better dev ranks help only slightly, and even the top decile averages -0.09 in validation.

6.5 Heatmaps — robustness across neighbours

At $H = 2$, reversion, $m = 1$, development Sharpe is positive in nearly every ($L \times k$) cell for every factor set and policy (≈ 0.1 – 0.9), so the day-2 reversal is a broad effect, not an isolated peak. $H = 1$ is weak or negative almost everywhere, and long holds ($H = 5$ – 10) under policy A are frequently negative.



Development net Sharpe over lookback $\times$ k (m = 1, H = 2, reversion); rows = factor sets, columns = policies; the primary is outlined.



Development net Sharpe over k x H (L = 20, m = 1, reversion).

6.6 Selection funnel (development)

	configs remaining
cumulative filter	
all selectable configs	10800
trades in 100% of months	2532
≥ 150 entries	2532
Sharpe > 0 in both halves	509
≥ 70% of years positive	183
positive ex top-1% days	10
policy ≠ D (stack)	10

The monthly-trade requirement removes every configuration with $k \geq 1.25$ (2,532 survivors all have $k \leq 1.0$); positive Sharpe in both halves removes most of the rest; the "ex top-1% days" test is the tightest. Removing the best 1% of days (≈ 22 days) turns the median configuration's annual return to -21% . The best near-misses — failing *only* the monthly-trade filter — score higher on the neighbourhood metric (up to 0.72, e.g. `GLD|L250|m2|k1.25|H1|A|rev`) than anything that passed. **No filter was loosened.**

6.7 Development shortlist and statistical checks

	nbhd median	dev Sharpe	entries	% yrs +	ex-top1% ret	DSR	placebo pct	sign- flip	boot lo	boot hi
config_id										
GLD L40 m1 k1 H3 A rev	0.431	0.707	398	0.889	0.003	0.020	99.300	-0.796	0.059	1.355
SPY+GLD L60 m1 k0.5 H10 B rev	0.415	0.721	717	0.889	0.034	0.021	99.600	-0.865	0.097	1.378
GLD L20 m1 k0.75 H2 A rev	0.408	0.829	655	0.889	0.033	0.041	99.600	-0.962	0.179	1.502
SPY+GLD+TLT L120 m1 k0.75 H3 B rev	0.401	0.778	659	0.778	0.042	0.033	99.700	-0.912	0.037	1.575
SPY+GLD L60 m1 k0.5 H3 B rev	0.400	0.631	829	0.778	0.002	0.011	98.400	-0.793	0.019	1.269
SPY+GLD L60 m1 k0.5 H5 B rev	0.365	0.740	734	0.778	0.041	0.024	98.700	-0.886	0.125	1.383
SPY+GLD+TLT L120 m1 k0.75 H5 B rev	0.352	0.647	571	0.778	0.001	0.013	98.700	-0.761	-0.100	1.442
SPY+GLD L250 m1 k0.5 H2 A rev	0.332	0.667	775	0.778	0.005	0.015	98.500	-0.798	-0.023	1.381
GLD L20 m1 k0.75 H3 A rev	0.311	0.795	504	0.889	0.051	0.035	99.700	-0.887	0.074	1.477
SPY+GLD L20 m2 k0.5 H10 A rev	-0.192	0.691	204	0.778	0.016	0.017	98.000	-0.721	0.029	1.394

All ten use the imposed reversion direction and nine use $m = 1$. Every one beats 98–99.7% of random-timing placebos and loses when inverted, and most bootstrap intervals exclude zero — but every Deflated Sharpe Ratio is 0.01–0.04. The DSR benchmark SR_0 is ≈ 1.41 (annualised): with 10,800 configurations whose Sharpe ratios spread widely, a top development Sharpe of 0.6–0.8 is what the best of many noise strategies would show. The configurations are highly correlated, so the effective number of independent trials is smaller than 10,800 and the DSR is conservative — but not by enough to change the conclusion. The tenth shortlisted configuration has a negative neighbourhood median (−0.19) and entered only because exactly ten passed.

6.8 Validation (evaluated once, logged)

	val Sharpe	entries	% months ≥ 1	% yrs +	Sharpe H1	Sharpe H2	ex-top1% ret	failed filters
config_id								
GLD L20 m1 k0.75 H2 A rev	0.459	384	1.000	0.800	0.209	0.629	-0.088	ex_top_days
SPY+GLD L60 m1 k0.5 H3 B rev	0.173	417	1.000	0.800	-0.032	0.305	-0.189	halves, ex_top_days
SPY+GLD+TLT L120 m1 k0.75 H5 B rev	0.239	284	0.983	0.800	0.529	0.076	-0.167	months, ex_top_days
SPY+GLD+TLT L120 m1 k0.75 H3 B rev	0.277	342	0.983	0.600	0.610	0.098	-0.159	months, years, ex_top_days
SPY+GLD L250 m1 k0.5 H2 A rev	0.387	439	1.000	0.600	-0.170	0.727	-0.122	halves, years, ex_top_days
SPY+GLD L60 m1 k0.5 H10 B rev	0.086	355	0.967	0.800	-0.090	0.194	-0.217	months, halves, ex_top_days
SPY+GLD L60 m1 k0.5 H5 B rev	0.172	367	0.983	0.800	-0.025	0.299	-0.188	months, halves, ex_top_days
GLD L20 m1 k0.75 H3 A rev	0.199	295	1.000	0.400	-0.290	0.491	-0.167	halves, years, ex_top_days
SPY+GLD L20 m2 k0.5 H10 A rev	-0.381	118	1.000	0.200	-0.838	-0.138	-0.351	halves, years, ex_top_days
GLD L40 m1 k1 H3 A rev	0.448	232	0.967	0.600	-0.901	1.192	-0.112	months, halves, years, ex_top_days

No shortlisted configuration passed the hard filters on validation. All ten fail the "ex top-1% days" test (without the best ≈ 13 days they lose 9–35% per year); several also fail positive Sharpe in both halves or the monthly-trade rule.

6.9 The frozen configurations

Primary GLD|L20|m1|k0.75|H2|A|rev ; alternates SPY+GLD|L60|m1|k0.5|H3|B|rev and SPY+GLD+TLT|L120|m1|k0.75|H5|B|rev (chosen by fewest validation failures, then combined validation-Sharpe and neighbourhood rank).


config	GLD L20 m1 k0.75 H2 A rev		SPY+GLD L60 m1 k0.5 H3 B rev		SPY+GLD+TLT L120 m1 k0.75 H5 B rev	
period	development	validation	development	validation	development	validation
net Sharpe	0.829	0.459	0.631	0.173	0.647	0.239
gross Sharpe	0.896	0.539	0.712	0.264	0.704	0.306
ann. return	0.302	0.131	0.278	0.058	0.289	0.078
ann. vol	0.365	0.285	0.440	0.335	0.446	0.329
max DD	-0.687	-0.575	-0.730	-0.849	-0.986	-0.401
NW t	2.409	0.928	1.862	0.384	1.731	0.523
trade hit rate	0.527	0.518	0.569	0.568	0.574	0.567
profit factor	1.299	1.176	1.203	1.058	1.260	1.099
entries	655.000	384.000	829.000	417.000	571.000	284.000
entries/month	6.359	6.400	8.049	6.950	5.544	4.733
% days in mkt	0.604	0.611	0.925	0.912	0.906	0.869
% yrs positive	0.889	0.800	0.778	0.800	0.778	0.800
Sharpe 1st half	0.898	0.209	0.510	-0.032	0.678	0.529
Sharpe 2nd half	0.755	0.629	0.763	0.305	0.614	0.076
ann. ret ex top1% days	0.033	-0.088	0.002	-0.189	0.001	-0.167
corr GDX	0.098	0.134	-0.010	0.012	-0.049	0.033

Statistical checks (development):

	GLD L20 m1 k0.75 H2 A rev	SPY+GLD L60 m1 k0.5 H3 B rev	SPY+GLD+TLT L120 m1 k0.75 H5 B rev
dev net Sharpe	0.829	0.631	0.647
Deflated Sharpe (N=10,800)	0.041	0.011	0.013
SR0 benchmark (ann.)	1.408	1.408	1.408
Prob. Sharpe > 0	0.994	0.968	0.971
placebo percentile	99.600	98.400	98.700
placebo 95th pct Sharpe	0.484	0.456	0.507
sign-flip Sharpe	-0.962	-0.793	-0.761
bootstrap 95% CI low	0.179	0.019	-0.100
bootstrap 95% CI high	1.502	1.269	1.442
skew	1.097	0.296	0.069
kurtosis	17.155	6.692	9.929

The primary trades about 6.4 times a month (655 dev / 384 validation entries), is in the market $\approx$ 60% of days, holds 2 days on average, turns over $\approx$ 115–120 units a year, and is nearly uncorrelated with GDX (corr 0.10–0.13), SPY and GLD. Its main weaknesses are its return distribution (kurtosis $\approx$ 17: the P&L depends on a few large days) and the validation decline (NW t-stat 2.41 $\rightarrow$ 0.93).


6.10 Equity curves, monthly entries and yearly returns



Cumulative net P&L (1x notional, additive) and drawdowns vs buy-and-hold GDX excess return.

	PRIMARY	Alt 1	Alt 2	Buy&hold GDX
year				
2008	0.494	0.912	1.021	-0.107
2009	0.279	0.093	0.222	0.443
2010	0.235	-0.108	-0.234	0.336
2011	0.218	0.109	0.314	-0.119
2012	0.268	0.131	0.136	-0.050
2013	-0.143	0.491	0.309	-0.672
2014	0.667	0.411	0.586	-0.059
2015	0.342	-0.066	-0.029	-0.188
2016	0.242	0.415	0.155	0.546
2017	-0.072	0.168	0.073	0.137
2018	0.189	0.006	0.242	-0.076
2019	0.023	-0.357	-0.124	0.350
2020	0.186	0.149	0.092	0.373
2021	0.328	0.324	0.109	-0.061
dev ann. Sharpe	0.829	0.631	0.647	0.032
val ann. Sharpe	0.459	0.173	0.239	0.413

Yearly net returns (sums of daily returns, 1x notional). The primary had 12 of 14 years positive; validation years: 2017 -7.2%, 2018 +18.9%, 2019 +2.3%, 2020 +18.6%, 2021 +32.8%. Buy-and-hold GDX has a validation Sharpe of 0.41 — close to the primary's 0.46.



New entries per calendar month. The primary and alternate 1 trade every month in both periods; alternate 2 misses one month in 2017.

6.11 Does it depend on 2008 / 2020? Does it survive costs?

	all years		ex 2008		ex 2020		ex 2008 & 2020	
	Sharpe	ann. return	Sharpe	ann. return	Sharpe	ann. return	Sharpe	ann. return
PRIMARY	0.709	0.239	0.695	0.212	0.749	0.244	0.745	0.215
Alt 1	0.487	0.197	0.364	0.136	0.512	0.201	0.379	0.135
Alt 2	0.519	0.211	0.389	0.143	0.558	0.221	0.420	0.147

The primary does **not** depend on 2008 or 2020 (Sharpe 0.71 all years; 0.70 ex 2008; 0.75 ex 2020; 0.75 ex both). The two alternates lean on 2008 ($\approx 0.5 \rightarrow \approx 0.37$ without it).

	development	validation
gross (0 bps)	0.90	0.54
2 bps	0.83	0.46
5 bps	0.73	0.34

The edge survives 5 bps per unit traded (validation Sharpe 0.34), but each basis point removes a visible share, and with ≈ 115 –120 units of turnover a year, costs of 2 bps take $\approx 2.4\%$ a year. Costs are not the binding problem; the thinness and tail-dependence of the edge are.

6.12 Walk-forward of the selection rule

	config chosen (data to Y-1)	train Sharpe	test-year Sharpe	entries
test year				
2011	GLD L250 m1 k0.5 H2 A rev	1.18	1.48	99
2012	GLD L250 m1 k0.5 H2 A rev	1.21	-0.58	81
2013	GLD L250 m1 k0.5 H3 A rev	1.09	-1.14	70
2014	GLD L250 m1 k0.5 H2 A rev	0.87	0.08	85
2015	GLD L120 m2 k0.5 H2 A rev	0.87	-0.51	86
2016	SPY+GLD L120 m1 k0.75 H3 B rev	0.74	0.16	71
2017	GLD L40 m1 k1 H3 A rev	0.71	-1.78	42
2018	GLD L20 m1 k0.75 H3 A rev	0.72	-0.35	58
2019	GLD L20 m1 k0.75 H2 A rev	0.75	0.11	69
2020	SPY+GLD+TLT L250 m2 k1 H1 A rev	0.97	-0.85	49
2021	SPY+GLD L60 m1 k0.5 H3 B rev	0.45	1.20	77

9 distinct configurations in 11 years, 2.7 of 7 parameters changing per switch (red flag raised); only the reversion direction was stable. Test-year Sharpe was negative in 6 of 11 years despite training Sharpe of 0.45–1.2 every year; stitched out-of-sample Sharpe **-0.19** (-6.3% a year). The primary itself was the walk-forward choice in only 1 year (2019).

7. Out-of-sample runner (ready, not yet run on real OOS data)

```
python scripts/run_oos.py --is data/raw/gold_etf_factor_regression.xlsx --oos data/raw/<OOS file>.xlsx \
--config results/frozen_config.json --out results/oos/
```

- Refuses to run if `frozen_config.json`'s SHA-256 differs from the hash recorded at freeze, or if the IS workbook is not the one the config was frozen on.
- Concatenates IS and OOS prices with seam checks: identical overlapping rows are dropped, differing overlaps or a gap of more than 4 weekdays are refused.
- Uses IS history as warm-up so rolling betas and $\hat{y}$ are live on OOS day 1 and keep updating (part of the frozen rule); no grid parameter changes. Frozen configs start flat and may first trade at the last IS close.
- Reports OOS-only metrics, hard-filter pass/fail, Sharpe at 0/2/5 bps, with/without 2008 and 2020, all Step-10 diagnostics, and an IS-vs-OOS side-by-side table; the run manifest records a check that the frozen development Sharpe still reproduces (currently within 1e-6).

8. Interpretation

- **There is information in the residual** — placebo percentiles of 98–99.7%, consistently losing sign-flips, reversion beating the estimated direction in both periods, and a broad positive region in the heatmaps all say the residual's timing is not random.
- **But it is small relative to the search** — the DSR, the weak dev→validation rank correlation (0.20), the negative walk-forward and the tail dependence (kurtosis 17; losses once the top 1% of days are removed) all say the *selected* configuration's Sharpe is mostly selection luck on top of a modest true effect.
- **The model structure is mis-aimed** — the effect is a day-2 reversal while Stage 2 forecasts day 1; the estimated-direction configurations therefore underperform and the chosen ones rely on an imposed direction that Stage 2 does not endorse.
- **The validation period was hard for this whole family** (grid median -0.27 vs buy-and-hold GDX $+0.41$), so part of the failure is regime, not only overfitting.

9. Next steps and possible changes

Any change below is made **after** seeing validation, so it must be logged as a protocol deviation and judged only on data that has not been used — the coming OOS file (or a later period). The frozen configuration should stay untouched and be evaluated on OOS first, as the honest test of what was built.

9.1 Immediately

1. **Run the frozen primary on OOS unchanged** with `run_oos.py`; report it against its flag (it failed the validation filters) and against buy-and-hold GDX. Treat it as a test, not a deployment.
2. **Record the frozen hash externally** (`6981dc95...`) so the OOS run can be checked with `--expected-sha256`.

9.2 Model changes suggested by the evidence (each needs fresh data to judge)

1. **Fix the horizon mismatch.** Make Stage 2 forecast the 2-day return ($\text{ExRet}_{t+1} + \text{ExRet}_{t+2}$) or regress on lagged residuals (e_t, e_{t-1}) jointly, so the estimated direction can capture the day-2 reversal instead of relying on an imposed sign.
2. **Use position sizing instead of a binary gate.** Scale exposure by $|\hat{f}|$ / forecast volatility (or target a fixed volatility), capped at 1x. This addresses the kurtosis / tail-day dependence that failed the "ex top-1% days" filter more directly than tuning k .
3. **Guarantee monthly activity differently.** The strongest near-misses ($k = 1.25, L = 250$) fail only the monthly rule. A quantile-based threshold (trade the top $q\%$ of $|z|$ within a rolling month) or a "at least one trade per month, taking the largest $|z|$ " fallback would keep the requirement while allowing a higher-conviction threshold.
4. **Ensemble instead of a single winner.** Average positions across a pre-declared neighbourhood (e.g. all $L \times k$ cells at $H = 2$, reversion, $m = 1$) rather than picking one cell; this directly targets the walk-forward instability.
5. **Execution realism.** Signal and execution at the same close is optimistic (market-on-close orders are submitted before the close). Test signal at $\approx 15:50$ prices or execution at the next close; the day-2 location of the edge suggests a one-day delay may cost less than expected. Add short-borrow cost and a slippage model beyond 2 bps.
6. **Richer residual model.** Consider miners-specific drivers (oil/energy as a cost proxy, USD, equity-sector factors) or GDXJ-GDX spread information; test whether the residual reversal is stronger after large GLD moves or on high-volume days.

9.3 Protocol improvements

1. **Smaller, pre-registered grid.** 10,800 trials made the DSR hurdle ≈ 1.4 annualised Sharpe. Fixing dimensions that the evidence already settles (direction = reversion, $H \approx 2$) and searching only a few others would lower the hurdle legitimately.
2. **Effective number of trials.** Estimate it from the correlation structure of the 10,800 return series (e.g. clustering or eigenvalue count) and report DSR at that N alongside $N = 10,800$.
3. **Longer holdouts.** With a thin effect, 5 years of validation cannot separate Sharpe 0.3 from 0; plan the next evaluation to use the full OOS period plus any later data, with nothing re-tuned on it.
4. **Cluster-robust event-study bands.** The current ± 1 s.e. bands treat overlapping events as independent and are too narrow.

10. Reproducing everything

```
cd gdx_strategy
pip install -r requirements.txt
python -m pytest -q                # 83 tests
python scripts/run_is.py           # development stage only (no validation access)
python scripts/run_is.py --report  # rebuild report.html / grid_results.parquet from the frozen config
python scripts/run_oos.py --is data/raw/gold_etf_factor_regression.xlsx --oos data/raw/<OOS>.xlsx --out results/oos/
```

Key outputs in `results/`: `frozen_config.json` (+ `.sha256`), `report.html`, `dev_shortlist.csv`, `validation_shortlist.csv`, `walk_forward.csv`, `all_configs_summary.csv.gz` (every configuration $\times$ period), `validation_log.jsonl` (every validation access, including the post-freeze reporting entry), `figures/`.